

# THU SUMMER SCHOOL 2025

CHANTAL DEUSCH | JEREMY DIEM | JAN GASCHLER | SERHAT GÜREL | PAULINA PYCZOT | VALENTIN TALMON-L'ARMÉE

## Image Detection with YOLO AI on a Raspberry Pi




# Agenda

1. Setup Raspberry Pi
2. Licence Plate Recognition
3. Face Recognition



# Visual Process Structure



Raspberry Pi

Taking Pictures



Detect Plate



Reading Plate



Live Stream



Detect Person  
and Face



Learn Face



# Set Up Raspberry Pi

To set up the Raspberry Pi, you will need the following physical components:

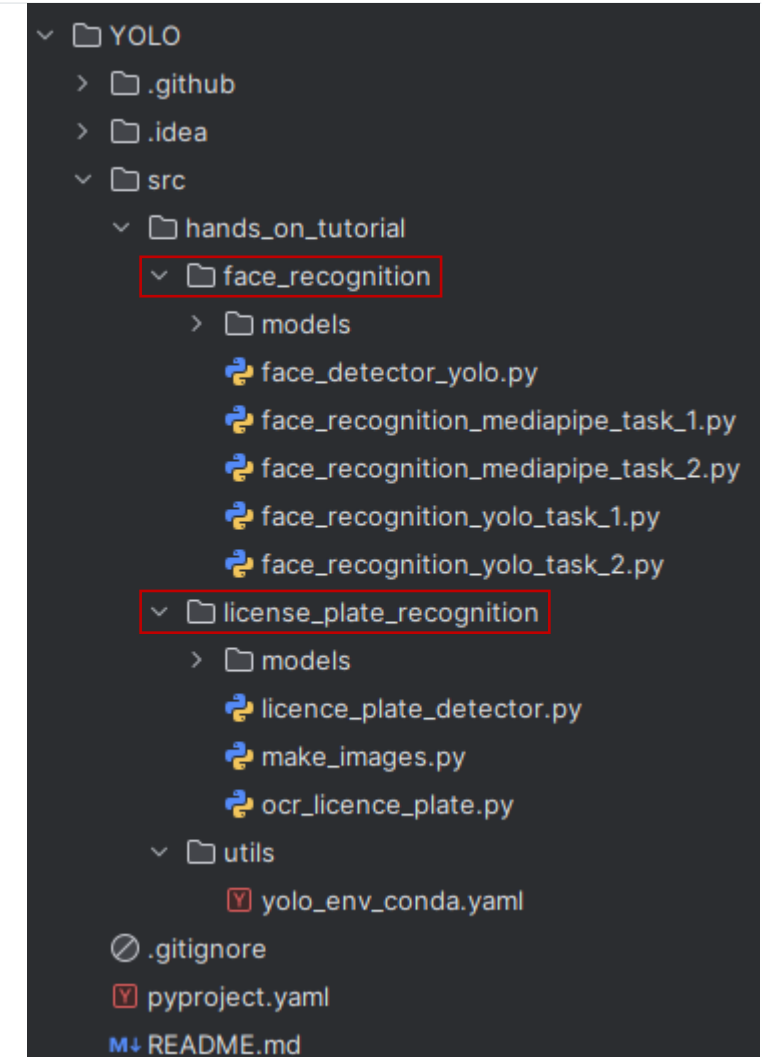
- Pre-configured SD card
- Mouse
- Keyboard
- Monitor with HDMI support
- USB-compatible camera
- Power supply & Micro-HDMI to HDMI cable



# Set Up Raspberry Pi

## Setting Up the Project

1. Clone Git repository
  - `git clone https://github.com/TOPR-yoloteam/YOLO.git`
2. Navigate into the project folder
  - `cd ./YOLO`



# Set Up Raspberry Pi

## Setting Up the Conda Environment

### Why do we need a virtual environment?

- › Keeps dependencies organized
- › Avoids version conflicts
- › Isolates the project from global Python
- › Makes setup easier and cleaner

#### Nice to know:

To activate your conda environment, use  
`$ conda activate <env>`

To deactivate an active environment, use  
`$ conda deactivate`

### Installation of Conda

1. Download the Miniconda installer and make it executable
  - `wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-aarch64.sh`
  - `chmod +x Miniconda3-latest-Linux-aarch64.sh`
2. Run the installer
  - `./Miniconda3-latest-Linux-aarch64.sh`
  - Press ENTER, scroll to the end of the EULA (= End User License Agreement)
  - Press `q` and type 'yes' to accept the terms
  - Press ENTER to install Miniconda3 into `/home/admin/miniconda3`
  - Type 'yes' to activate conda on startup
3. Test it: Activate conda
  - `source ~/miniconda3/bin/activate`
  - You're now on your base environment!
  - Deactivate the environment with `conda deactivate`

# Set Up Raspberry Pi

## Installation of Libraries with .yaml

### Why .yaml?

- › Ensures consistent installation of required packages
- › Makes it easy to recreate and share the exact environment setup
- › All needed dependencies (general, license plate recognition & face recognition) in one place

### Setting Up the virtual Conda Environment with .yaml

1. Optional: Create a project folder and navigate into it
  - `mkdir <folder_name>`
  - `cd <folder_name>`
2. Create a new conda environment
  - `conda env create -f ./YOLO.yaml`
3. Activate the environment
  - `conda activate yolo_env_conda`
  - The location of your environment is `/home/admin/miniconda3/envs/yolo_env_conda`

```
hands_on_tutorial\...\yolo_env_conda.yaml x
1  name: yolo_env_conda
2  channels:
3    - conda-forge
4  dependencies:
5    # General libraries
6    - python=3.9          # Python version
7    - pip                 # Python package manager
8    - numpy=1.26.4        # Numerical operations and array handling
9    - matplotlib          # Visualization and plotting
10   - opencv               # Computer vision library (image and video processing)
11   - pygame               # Multimedia handling (e.g., simple UI, window management)
12
13   # Special libraries for text, characters, and license plate recognition
14   - pytesseract          # Python binding for Tesseract OCR
15   - tesseract            # Tesseract OCR engine (backend for pytesseract)
16
17   # Special libraries for face/person recognition
18   - ultralytics          # YOLO models for object detection (including person detection)
19   - dlib                 # Machine learning toolkit (especially for facial landmarks)
20   - face_recognition      # High-level face detection and recognition built on dlib
21
22   # pip-only packages
23   - pip:
24     - mediapipe           # Google framework for face/hand/pose detection
25     - torch==1.10.0       # PyTorch framework (deep learning backend)
26     - tensorflow==2.8.0   # TensorFlow framework (for models that require it)
27
28   # Create environment with: conda env create -f yolo_env_conda.yaml
29   # Activate environment with: conda activate yolo_env_conda
```

🔧 If there should be a problem using the .yaml file you can setup your conda environment manually by using the installation file `YOLO/src/hands_on_tutorial/utils/YOLO_man.txt`

# Setup Raspberry Pi

## Setting Up the Python Development Environment

### Thonny

1. A simple and beginner-friendly Python development environment
2. Pre-installed on Raspberry Pi OS
3. Built-in debugger

### How to Start and Use Thonny

1. Type *thonny* in the Terminal
2. Use the file explorer to navigate to the Python file you want to edit by using the *Load* button
3. To run the script, click the green *Run* button at the top (or press F5 on your Keyboard)



# License Plate Recognition



# License Plate Recognition

## Helpful tips

- › To see changes better, use the following code after changing an image
- › Code to be written is marked with #TODO X in the source code
- › Google loves you, ChatGPT doesn't <3

```
cv2.imshow( winname: "image", image)  
cv2.waitKey(0)
```

```
#TODO 1
```

## Taking Pictures

- › To recognize license plates, we need images. These images must be captured first.  
We do this using the *Image* class in the *make\_images.py* file.
  - Run the script to start capturing images from the webcam



# License Plate Recognition

## Detect License Plate

- › In order to be able to read the text from license plates later, we must first recognize the license plate as such
- › **Task 1: Modify the Python code to detect a license plate and draw its corresponding bounding box**
- › *Hint:* The Ultralytics docs homepage provides information on how to use a YOLO model :')



# License Plate Recognition

## Detect License Plate

- › To achieve the best OCR results, we now want to save the license plate as a single image
- › **Task 2: Save the license plate as a single image in a separate folder named "license\_plates"**
- › *Hint:* Use the plates bounding box coordinates!





# License Plate Recognition

## Read License Plate

- › Colored images can cause problems in character recognition, therefore we need to preprocess the input image
- › **Task 1: Preprocess Image on color, threshold, dilation and contours**
- › *Hint:* Use the OpenCV library for image preprocessing. For contours use the given method "find\_and\_sort\_contours"



## Read License Plate



- › If all steps are done correctly, start with the OCR!
- › **Task 2:** Extract text from the contours and print the results
- › *Hint:* What does "extract\_text\_from\_contours" return?



```
Image: image_0_0.png
```

```
Text: ULDT805 | Probabilities: 86.33
```

# Face Recognition





Please choose between:


- YOLO
  - › In this tutorial, we explore how the YOLOv8n-face model can be used to detect human faces in real time
  - › The YOLOv8n-face architecture is fine-tuned specifically for faces, combining high detection accuracy with a relatively small model size (only 1 object class)
  - › Although **YOLO** is highly optimized for GPU acceleration and real-time performance, it tends to lag more than **MediaPipe**, especially on resource-constrained devices like a Raspberry Pi
- MediaPipe
  - › A framework by Google providing lightweight models, particularly FaceMesh for detailed facial landmark detection.
  - › Often highly performant, even on resource-constrained devices like the Raspberry Pi.
  - › Focuses on extracting facial features (landmarks) which can then be used for recognition by comparing their relative positions.

If you want to solve the MediaPipe exercises, please skip to page **33**

## YOLO

### Task 1:

- › Implement **face detection** based on YOLO.
- › Detect **faces** and draw bounding boxes around them.



## YOLO Task 1 tips

### Task 1.1:

- › Ensure the `self.model_path` correctly points to the "yolov8n-face.pt" file within the project structure.

### Task 1.2:

- › You run detection by calling the model object directly: `results = self.model(...)`.
- › Pass the video frame as the main argument to the model.
- › **Crucial:** Use the `classes=[0]` argument to tell the model to *only* detect faces (which is class 0 in this specific face model).

### Task 1.3:

- › The results object contains detections. You often need to loop through results, then `result.bboxes`, then each box.
- › The coordinates are typically stored within a box object, often under the attribute `.xyxy[0]`.
- › The coordinates might be floating-point numbers. Use `map(int, ...)` to convert them to integers before drawing.
  - › *Example:* `x1, y1, x2, y2 = map(int, box.xyxy[0])`

## YOLO Task 1 tips

### Task 1.4:


- › You can use OpenCV's drawing function: `cv2.rectangle(...)`.
- › Remember the required arguments: `cv2.rectangle(image, pt1, pt2, color, thickness)`.
  - › *pt1* is the top-left corner (x1, y1).
  - › *pt2* is the bottom-right corner (x2, y2).
  - › *Color* is a BGR tuple, e.g., (0, 255, 0) for green.
  - › *Thickness* is an integer, e.g., 2.

If you want to solve the MediaPipe exercises, please skip to page **33**

## YOLO

### Task 2:

- › Extend the system by adding face recognition.  
Person detection is **already provided**.  
Implement recognizing known faces and mark them in the image.



## YOLO Task 2 tips

### Task 2.1:

- › Make sure you're reading image files from the correct directory (`self.faces_dir`).
- › Use `face_recognition.load_image_file()` to load each image.
- › Get the face encoding using `face_recognition.face_encodings(loaded_image)[0]`.
  - › Note the `[0]` assumes one face per known image file.
- › Store the encoding in `self.known_face_encodings` and the name (from filename) in `self.known_face_names`. Keep these lists synchronized!

### Task 2.2:

- › Get the `face_location` passed to the `save_face` function.
- › Crop the face from the stored `self.current_frame` using these coordinates: `face_image = self.current_frame[top:bottom, left:right]`.
- › Consider expanding the coordinates slightly before cropping to get more context, but be careful not to go outside the frame boundaries (`max(0, ...)`, `min(width, ...)`).
- › Save the `face_image` using `cv2.imwrite()` into `self.faces_dir`. Give it a filename based on the entered name.
- › **Very Important:** After saving, immediately call `self.load_known_faces()` again so the system recognizes the newly added person right away.

## YOLO Task 2 tips

### Task 2.3, 2.5., 2.6, 2.7:

- › Use the self.state variable ("normal" or "entering\_name") to control behavior.
- › In mouse\_callback: If state is "normal", check if the click (x, y) is inside any stored button area. If yes, change self.state to "entering\_name" and store which face was clicked (self.selected\_face\_loc).
- › In the main loop (run): If self.state is "entering\_name", handle keyboard input (cv2.waitKey(1)):key == 13 (Enter): Save the face if self.current\_text is not empty, then switch back to "normal" state.
  - › key == 27 (Esc): Switch back to "normal" state.
  - › key == 8 (Backspace): Remove last character from self.current\_text.
  - › Printable characters: Append chr(key) to self.current\_text.
- › Draw the text input UI (draw\_text\_input) and highlight the selected face only when in the "entering\_name" state.

## YOLO Task 2 tips

### Task 2.4:

- › You need face locations and encodings for the *current* video frame first.
- › Use `face_recognition.face_locations()` and `face_recognition.face_encodings()`.
- › **Remember:** Convert the OpenCV frame (BGR) to RGB before passing it to face\_recognition functions: `rgb_frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)`.
- › Use `face_recognition.compare_faces(self.known_face_encodings, encoding_to_check)` to get potential matches.
- › Use `face_recognition.face_distance(self.known_face_encodings, encoding_to_check)` to find the distance for each comparison. Lower distance means a closer match.
- › Find the best match using `np.argmin(face_distances)` on the distances array.
- › Verify the best match: Check if `matches[best_match_index]` is True.



## YOLO Task 2 tips

### Task 2.4:


- › Draw the name label using `cv2.putText()`.
- › Draw the "Learn Face" button (using `cv2.rectangle` and `cv2.putText`) **only** if the face is "Unknown" and the system is in the "normal" state.
- › For the button click to work later, store the button's screen coordinates *and* the corresponding face location (top, right, bottom, left) together (e.g., in `self.button_area`).

If you want to solve the YOLO exercises, please go back to page 25

## MediaPipe

### Task 1:

- › Implement **face detection** based on MediaPipe.
- › Face recognition is not required in this task.
- › Detect **faces** and draw bounding boxes around them.



## MediaPipe Task 1 tips

### Task 1.1:

- › Create the FaceMesh instance using `mp.solutions.face_mesh.FaceMesh(...)`.
- › You can set parameters like `max_num_faces=...` and `min_detection_confidence=...` during initialization (check MediaPipe documentation for details).

### Task 1.2:

- › **Important:** MediaPipe expects images in RGB format, but OpenCV captures in BGR. Convert the frame: `image_rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)`.
- › Call the `process()` method on your `self.face_mesh` object, passing the `image_rgb`.

### Task 1.3:


- › Use OpenCV's `cv2.rectangle(image, pt1, pt2, color, thickness)`. `pt1` should be the top-left corner: (left, top). `pt2` should be the bottom-right corner: (right, bottom).
- › Choose a color (BGR tuple like (0, 255, 0) for green) and thickness.

If you want to solve the YOLO exercises, please go back to page 25

## MediaPipe

### Task 2:

- › Extend the system by adding face recognition.
- › Person detection is **already provided**.
- › Implement recognizing known faces and mark them in the image.



## MediaPipe Task 2 tips

### Task 2.1:

- › You don't need all 400+ landmarks from FaceMesh. Use the provided `key_landmarks_indices` list to select a smaller, specific set (e.g., 50 landmarks).
- › Loop through `key_landmarks_indices`. For each index `idx`, get the landmark data: `landmark = face_landmarks.landmark[idx]`.
- › Store the **normalized** coordinates `landmark.x` and `landmark.y` (values between 0.0 and 1.0).
- › Create a flat NumPy array from these coordinates (e.g., `[x1, y1, x2, y2, ..., x50, y50]`). The size must be consistent!

### Task 2.2:

- › This happens in the `compare_landmarks` function.
- › You need to compare the detected landmarks array to *every* stored landmark sample for *every* known person (for `known_landmarks` in `person_landmarks_list`).
- › A good way to measure similarity is the Euclidean distance: `distance = np.linalg.norm(landmarks - known_landmarks)`.
- › Keep track of the `min_distance` found and the `best_match_index` (which person it corresponds to).

## MediaPipe Task 2 tips

### Task 2.3:

- › After checking all known landmarks, take the overall min\_distance.
- › Compare it to self.recognition\_threshold.
- › If min\_distance < self.recognition\_threshold, it's a match! Set is\_known\_face = True and get the name using best\_match\_index.
- › If the distance is larger, it's "Unknown" (is\_known\_face = False).
- › Consider calculating a confidence score based on how far below the threshold the distance is.